

# Welcome to Week 11 Lecture 2!

Data Science in Python &  
Machine Learning



# Neural Networks Review Poll:

1. Forward Propagation...
2. Backward Propagation...

# Learning Objectives

After this lesson you will be able to:

- ❑ Identify bias and variance in neural networks
- ❑ **Strategically** tune a neural network.

# Building & Tuning Neural Networks


# Cheatsheet: Neural Networks in Keras

(Copy or screenshot this slide for reference)

## A Keras model follows these main steps:

1. Instantiate the model type (Sequential for our models)
2. Add each layer in the order you want them to run
  - a. Each layer needs a **number of nodes** and an **activation function (hyperparameters)**
  - b. 1st layer must have defined input size (input\_dim=X\_train.shape[1])
  - c. Last layer:
    - i. Nodes = 1 for regression or binary classification, OR the number of target classes for multiclass classification.
    - ii. Activation must be appropriate to the problem.
      1. 'linear' for regression
      2. 'sigmoid' for binary classification
      3. 'softmax' for multiclass classification
3. Compile model
  - a. Loss appropriate for problem (MSE, BCE, or 'categorical\_crossentropy')
  - b. (optional but recommended) Optimizer ('adam' is good!)
  - c. (optional but recommended) additional metrics appropriate for the problem type. You can find [list of metrics here](#)
4. Fit model
  - a. Training set
  - b. Validation set
  - c. Epochs
  - d. (optional) Callbacks


# Tuning Neural Networks



[Image Source](#)

# Underfitting (High Bias)

A model that is **too simple** or **trained for too few epochs** won't learn to predict a dataset well.





# Attacking High Bias

**If model still seems to be learning:**

- Increase epochs
- Add nodes
- Add Layers

**If learning has leveled off:**  
increase model complexity



- Add nodes
- Add layers






# Attacking High Bias

**If Variance is low,  
attack the bias!**




# Overfitting (High Variance)

- High Variance: model performs worse on test data than on training data.



# Overfitting (High Variance)

- Neural networks are prone to overfitting
- **Attacking High Variance:**
  - Dropout Layers
  - Early Stopping
  - L1 and L2 Regularization
  - Decrease model complexity



# Dropout Layers

## What it does:

A dropout layer 'turns off' nodes at random in the layer before fit.


It turns off **different** nodes on each epoch.

## Why it works:


A network can become overly reliant on a few particular nodes which pick up very specific patterns.

Dropout forces a network to distribute the learning across many nodes and find more generalizable patterns

Image Source:  
Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting", JMLR 2014



(a) Standard Neural Net




(b) After applying dropout.

# Dropout Layers


How to do it:

```
1 # create and compile the model
2 from tensorflow.keras.layers import Dropout
3
4 model = keras.Sequential()
5 model.add(Dense(10, input_dim=X_train_processed.shape[1],
6                 activation = 'relu'))
7 model.add(Dropout(.2))
8 model.add(Dense(5, activation = 'relu'))
9 model.add(Dense(1, activation = 'sigmoid'))
10
11 model.compile(optimizer = 'adam', loss = 'bce', metrics = 'accuracy')
```

Percentage of nodes to 'drop' on each epoch



(a) Standard Neural Net



(b) After applying dropout.

# Callbacks: Early Stopping


## What it does:

Callbacks are functions that can be called during training. They check something about the training between epochs.

Early stopping stops training early.

## Why it works:

One source of overfitting is **over-training**. After a certain number of epochs a model will often begin to fit too tightly to training data. It learns TOO well. Early stopping stops the model before that can happen.




[Image Source](#)

# Callbacks: Early Stopping

```
from tensorflow.keras.callbacks import EarlyStopping
early_stopping = EarlyStopping(patience=3)
# train the model and save the history
history = model.fit(X_train_processed, y_train,
                    validation_data=(X_test_processed, y_test),
                    epochs=10,
                    callbacks=[early_stopping])
```

List of callbacks in fit call  
(There are many other  
callbacks you can use!)

If no reduction in  
validation loss in 3  
epochs, stop training



[Image Source](#)

# Advanced Technique: L1 and L2 Regularization

## What it does

- Neural network node = linear regression model
- L1 and L2 regularization to limit the change in weight values by applying a 'penalty' term.

## Why it works

- Weights can get too large
  - Model puts too much emphasis on a small subset of features
- Regularization L1 and L2 prevent this.



# Advanced Technique: L1 and L2 Regularization

How to do it:

```
from tensorflow.keras import layers
from tensorflow.keras import regularizers

layer = layers.Dense(
    units=64,
    kernel_regularizer=regularizers.l1_l2(l1=1e-5, l2=1e-4),
    bias_regularizer=regularizers.l2(1e-4),
    activity_regularizer=regularizers.l2(1e-5)
)
```

Image Source: [Keras Documentation](#)

# Reduce Model Complexity

**What it does:**

Reduce the number of nodes and/or layers in your model.

**Why it works:**

One source of overfitting is a model that is too complex. An overly complex model will fit to the noise in the data rather than the true signal.

**How it works:**

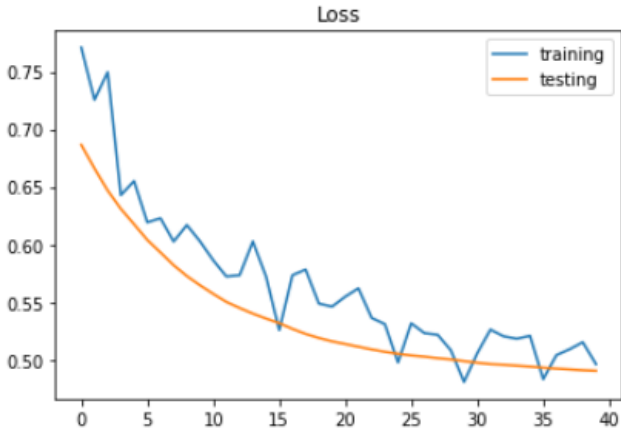
Create a model with fewer nodes and/or layers.

**(Remember to just change one thing at a time when experimenting)**

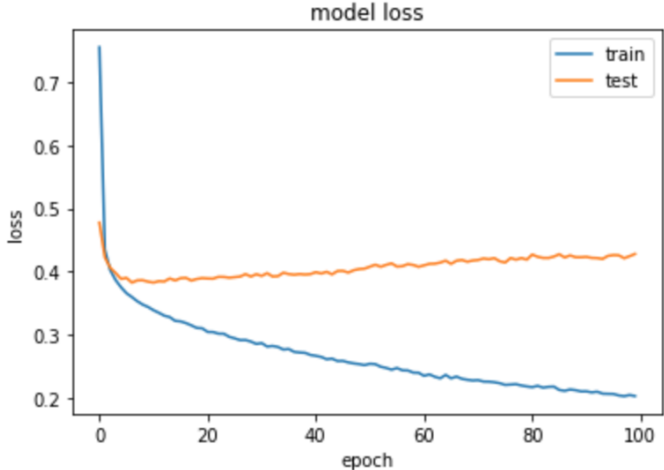
**Remember: Always evaluate each model with MULTIPLE metrics.**

**The best model is the model with the best metrics on test data,**  
regardless of overfitting!

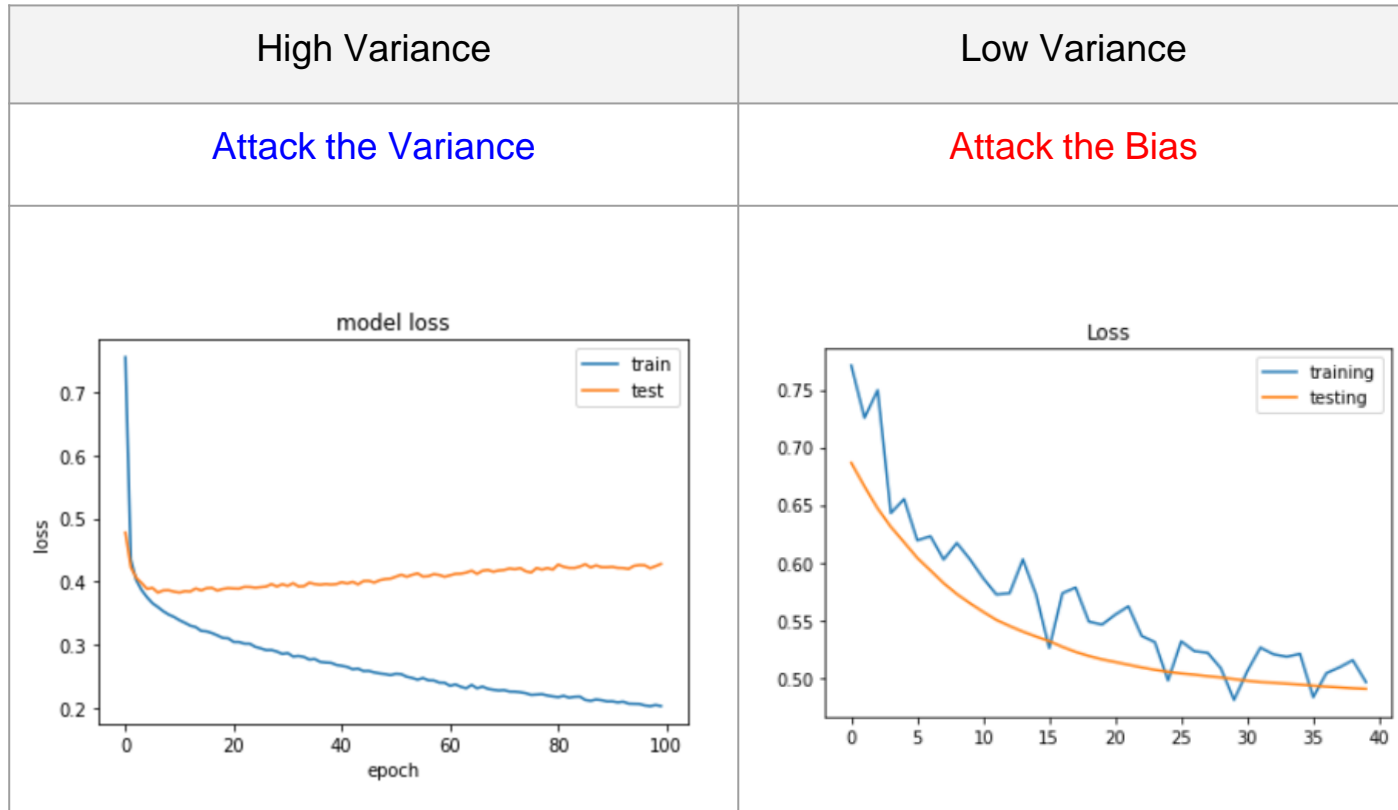
## Attack Bias or Attack Variance? (POLL)

Attack <b>Bias</b> or <b>Variance</b> ?	How?
<b>bias</b>	Increase complexity (nodes)
	Increase complexity (layers)
	Train longer

# Attack Bias or Attack Variance? (POLL)

Attack Bias or Variance?	How?
Variance	Reduce complexity (nodes and layers)
	dropout
	L1, l2 Early stopping

## Attack Bias or Attack Variance?



# Tuning a Model:

## **Strategies:**

1. More or fewer nodes
2. More or fewer layers
3. More or fewer epochs
4. Regularization
  - a. L1 or L2
  - b. Dropout layers
  - c. Early Stopping
5. Change activation functions
6. Change optimizers
7. And much more!

# Let's Tune a Model Together! (Code-along)

[Code-along Notebook](#)






# Your Challenge: Tune a Neural Network!

Your task will be to tune the classification model we created in the last lecture.

1. Make a copy of the notebook
2. Run the notebooks
3. Examine the learning plot and decide what you will try with the next model version.
4. (optional) copy/paste the model code from above and make changes to improve test accuracy

[Here is the starter notebook](#)



# Announcements: Belt Exam

- Belt exam code will be given to the eligible students before Friday evening
- To receive a code for this weekend: you must submit by Friday 9 am Pacific:
  - Week 9 (including all resubmits) and Week 10 assignments
- You will be given belt color based on
  - Black Belt: Your Score  $\geq 9.5$
  - Red Belt:  $8.0 \leq$  Your Score  $\leq 9.5$
  - Retake/Rollback: Your Score  $< 8.0$

**Everything covered in weeks 9 through 11 will be included in the belt exam, so you need to review it all.**

# Announcements: Presentations

- Presentations on Thursday next week.
- You will find out the order in class on Thursday, so everyone must be ready!
- You will have a max of 5 minutes (we will have a timer) so please practice!
- Should be aimed at a **Non-Technical Audience!**
  - What would a:
    - non-technical,
    - non-data scientist,
    - non-statistician
    - **NOT** understand?

TAs can give feedback on READMEs, slides, and they can be practice audiences if you ask them nicely!

# Assignments Due:

1. [Neural Network Exercise](#) - Optional Kaggle Competition
  - a. Explain your Model Changes
  - b. Recommend you submit an entry!
2. [Project 2 - Part 5](#) - Presentation slides
  - a. Remember: **NON-DATA SCIENCE AUDIENCE!!!**

## Daily Schedule

# Next Lecture: SQL Queries

Please Read:

1. [Intro to SQL](#)
2. [SQLAlchemy](#)
3. [SELECT and FROM](#)
4. [WHERE](#)
5. [Wildcards \(%\), and LIKE](#)
6. [Advanced WHERE](#)
7. [ORDER BY and LIMIT](#)

# Local Environment Installation

If you haven't already:  
Be sure you get this done before the end of the week

**DON'T WAIT UNTIL NEXT THURSDAY/FRIDAY**

Get help from instructors and TAs during the week on the [Environment Install Discord Channel](#)

Required assignment for next week AND  
**Critical** for first week of Data Enrichment